

Certifying Differential Invariants of Neural Networks using Abstract Duals

CHANDRA KANTH NAGESH, University of Colorado Boulder, USA

Neural networks are increasingly deployed in safety-critical domains, making the formal verification of their robustness against adversarial perturbations a paramount concern. While relational abstract domains like DeepPoly provide state-of-the-art certification, they suffer from the "wrapping effect," where overapproximation errors accumulate layer-by-layer, potentially failing to verify robust networks. In this work, we propose a complementary verification approach using *Abstract Dual Numbers* to compute global Lipschitz bounds via forward-mode automatic differentiation in the abstract domain. We contribute an algorithmic realization of this method that accounts for curvature error, a verification logic based on gradient bounds, and a "Gradient Instability" metric that serves as a precise predictor of verification failure. Our evaluation identifies a specific regime—shallow networks with smooth activations—where our global gradient method successfully certifies 16.1% of cases that DeepPoly fails to verify, highlighting the trade-off between local precision and global coherence.

1 Introduction

The deployment of deep neural networks in safety-critical domains such as autonomous driving and medical diagnosis has necessitated rigorous verification of their properties. One of the most critical properties is *local robustness*, which certifies that a network's prediction remains invariant under small perturbations to the input. DeepPoly[6] provides us with a clear method to certify output bounds for a given neural network. Let us consider the situation where we are given a trained neural network $f(\mathbf{I})$ on some set \mathbf{I} . Then, under some input $x_0 \in \mathbf{I}$ and perturbation ϵ , we need to verify that the output of f still satisfies the same result i.e. we want to verify that the neural network is invariant to a ϵ perturbation of the input. This can formally be represented as:

$$\forall x \in \mathbb{B}_\infty(x_0, \epsilon) : \text{argmax}(f(x)) = \text{argmax}(f(x_0))$$

where, \mathbb{B}_∞ is a ball of ϵ radius around the input provided by its L_∞ norm as shown in Fig 1. This is critical in a lot of scenarios as with many supervised learning problems, it is infeasible to verify the output of a neural network against all possible test sets. We want to formally prove that the model is invariant to these input changes.

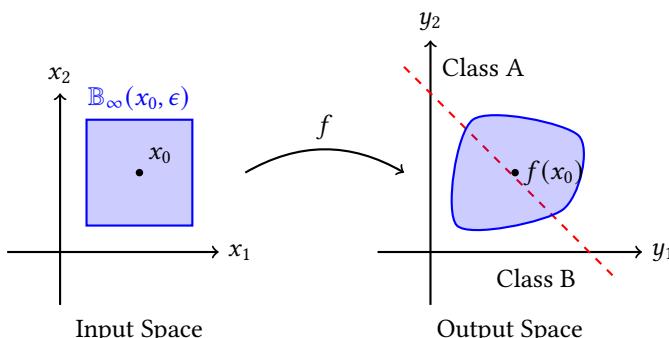


Fig. 1. The input region $\mathbb{B}_\infty(x_0, \epsilon)$ (left) is mapped by the neural network f to a region in the output space (right). Verification succeeds if the entire mapped region lies within the decision boundary (dotted red) of the correct class (Class A). This illustrates the robustness verification problem.

The earliest solutions that attempted to solve this problem are end-to-end/complete verification solutions. [4] introduced a tool called Reluplex, which is a solver based on SMT that combined the Simplex algorithm with ReLU activation functions to exactly verify small neural networks. Later, [7] used Mixed Integer Linear Programming (MILP) to write the problem as a commercial solution that exactly calculates robustness bounds. Although these solutions are sound and complete, they are NP-complete problems, known to not scale with the size of modern neural networks which are deployed in safety-critical applications.

For overcoming the scalability issue, Abstract Interpretation [1] and later for neural networks [2] has been adopted, which is less complete but more efficient because it over-approximates how the network behaves. [8] introduced a way to quickly compute robustness certificates with a linear bound propagation approach. The DeepPoly[6] approach pushed the boundaries even further by combining floating-point polyhedra with intervals, which, thanks to linearization, managed non-linearities such as Sigmoid and Tanh activations. More recent contributions, such as [5]

In our research, we ask the following question:

Can we continue to prove robustness properties of a neural network by leveraging the properties of dual numbers and information about the gradients?

We proceed with the foundation built on the *Abstract Dual Domain*. Our research is based on the theory of forward-mode automatic differentiation, which uses a concept called *dual numbers*[3]. We aim to analyze the sensitivity of the neural network to changes in the input, which can potentially provide a tighter bound on certain verification problems. In summary, our contributions are:

- *Algorithmic Realization of Abstract Duals*: We implement a sound abstract domain for dual numbers that explicitly handles curvature error for non-linear activations using Taylor expansion bounds.
- *Gradient-Based Verification Logic*: We formulate a sufficient condition for local robustness based on global Lipschitz bounds derived from the abstract gradient.
- *Gradient Instability Metric*: We introduce a diagnostic metric, $\mathcal{I}_{unstable}$, which detects zero-crossings in gradient intervals and serves as a predictor for the breakdown of local linearity.
- *Empirical Analysis*: We evaluate our method against DeepPoly, identifying a specific "shallow and smooth" regime where global gradient bounds outperform layer-wise relational abstractions.

2 Overview

We introduce the key insight behind our approach and define the Abstract Dual Domain in this section. Our approach is based on the following intuition: while relational abstractions such as DeepPoly[6] can be extremely precise, they often suffer from the “wrapping effect” where approximation errors accumulate layer by layer. In contrast, global gradient bounds can be sometimes a simpler yet effective verification domain, especially for shallow smooth networks.

2.1 From Relational Bounds to Gradient Analysis

The basic underlying concept of this work is to effectively bridge the gap between standard Abstract Interpretation and Forward-Mode Automatic Differentiation. Standard methods like DeepPoly track affine constraints ($x_j \geq \sum w_i x_i$) to bound the output values directly. However, precise reapproximating the feasible set at every non-linear layer introduces error which tends to accumulate with depth ($O(L)$). Our approach, instead, exploits the Mean Value Theorem to find

99 bounds on the output variation in terms of the gradient variation:

$$100 \quad |f(x) - f(x_0)| \leq \underbrace{\sup_{z \in X} \|\nabla f(z)\|}_{\text{Abstract Duals}} \cdot \|x - x_0\| \quad (1)$$

104 By over-approximating the sound of the gradient $\nabla f(z)$ for the entire input space X , we can
 105 guarantee robustness because the maximum possible change to the output will be too small to
 106 change the classification label. The key insight is that this “jumps” over the intermediate layers
 107 and thus avoids some of the wrapping errors from the layer-wise value propagation.

108 2.2 Abstract Dual Domain

110 To formalize this idea, we can start by lifting standard dual numbers to the domain of Affine
 111 Arithmetic. Consider the Dual domain $\hat{\mathcal{D}}$ as a product space of two affine forms representing the
 112 range of values and the range of gradients across a set. Then we can say:

113 *Definition 2.1.* An **Abstract Dual Number** $X \in \hat{\mathcal{D}}$ is a pair:

$$114 \quad X = \langle \hat{x}_{val}, \hat{x}_{grad} \rangle \quad (2)$$

117 where:

- 118 $\square \hat{x}_{val} = \alpha_0 + \sum_{i=1}^n \alpha_i \epsilon_i$ is the affine form representing the interval of neuron values.
- 119 $\square \hat{x}_{grad} = d_0 + \sum_{i=1}^n d_i \epsilon_i$ is the affine form representing the interval of partial derivatives with
 120 respect to the input.

121 2.3 Propagation Rules

123 For the dissemination of the abstract state X within a neural network, abstract transformers based
 124 on different types of layers are formulated.

125 **Linear Layers:** For a fully connected layer with a weight matrix W and a bias vector b , the
 126 transformation is calculated exactly thanks to the linearity of the dual algebra. The affine expressions
 127 for value and gradient are transformed into:

$$128 \quad Y = \langle W\hat{x}_{val} + b, W\hat{x}_{grad} \rangle \quad (3)$$

130 This is exact in the affine domain, which means that there is no additional error introduced here.

131 **Nonlinear Activation Functions:** For a smooth nonlinear activation function σ (for example,
 132 sigmoid), direct application to the affine expressions is not possible. The value part and the gradient
 133 part are handled separately.

134 *Value Component:* $\sigma(\hat{x}_{val})$ is approximated using a linear relaxation. We linearize the function
 135 around the center of the input affine form c . The new center is $\sigma(c)$, and the noise coefficients are
 136 scaled by the derivative $\sigma'(c)$. To ensure soundness, we add a linearization error term to the radius
 137 r , which depends on the maximum curvature (second derivative) of σ and the radius of the input
 138 interval.

$$139 \quad \hat{y}_{val} \approx \sigma(c) + \sigma'(c) \cdot (\hat{x}_{val} - c) + \epsilon_{err} \quad (4)$$

141 *Gradient Component:* By the chain rule, the gradient of the output is the product of the local
 142 derivative and the input gradient: $\nabla y = \sigma'(x) \cdot \nabla x$. In our abstract domain, we compute an affine
 143 approximation of the derivative $\hat{\sigma}'$ based on the output value \hat{y}_{val} . For the Sigmoid function, we
 144 use the property $\sigma'(x) = \sigma(x)(1 - \sigma(x))$ to approximate the derivative as $\hat{\sigma}' \approx \hat{y}_{val} \otimes (1 - \hat{y}_{val})$, where
 145 \otimes denotes affine multiplication. The output gradient is then obtained by multiplying this derivative
 146 approximation with the input gradient affine form given by $\hat{y}_{grad} = \hat{\sigma}' \cdot \hat{x}_{grad}$. This approach

preserves the correlations between the derivative and the gradient, providing a more precise abstraction than simple interval scaling.

2.4 Gradient Instability

A key insight from our analysis is the concept of *gradient instability*. We define a neuron n as having an unstable gradient if its gradient interval $[g, \bar{g}]$ contains 0:

$$\exists x_1, x_2 \in X_0 \text{ s.t. } \text{sign} \left(\frac{\partial f}{\partial n}(x_1) \right) \neq \text{sign} \left(\frac{\partial f}{\partial n}(x_2) \right) \quad (5)$$

Geometrically, this means that the function is non-monotonic w.r.t. neuron n over the input region. When this happens, the assumption of local linearity breaks down and linear relaxations—such as those in DeepPoly—become loose. Our approach explicitly tracks these gradient intervals, thus enabling us to identify whenever and wherever the network’s behaviour becomes hard to certify linearly.

3 Contributions

3.1 Algorithmic Realization of Abstract Duals

While the theoretical foundation rests on dual numbers, our practical contribution is the algorithmic realization of these concepts within the **Affine Arithmetic** domain. We implemented a custom OCaml module `AbstractDual` that handles the propagation of affine forms. Crucially, for non-linear activations σ , we implement a sound linearization that explicitly accounts for the curvature error. As seen in our implementation, the radius of the value component is expanded by a quadratic term derived from the Taylor expansion:

$$r_{new} = |\sigma'(c)| \cdot r_{old} + 0.5 \cdot \underbrace{\max_z |\sigma''(z)| \cdot r_{old}^2}_{\text{Linearization Error}} \quad (6)$$

This ensures that the abstract transformer remains sound even when the function has significant curvature, a detail often omitted in standard dual number formulations but critical for verification.

3.2 Gradient-Based Verification Logic

We formulate a robust verification condition based on the computed gradient bounds. Let $f_c(x)$ be the score of the correct class and $f_j(x)$ be the score of a competing class. The robustness margin is $M(x_0) = f_c(x_0) - f_j(x_0)$. To certify robustness over $\mathbb{B}_\infty(x_0, \epsilon)$, we compute the global Lipschitz constant K of the margin function $m(x) = f_c(x) - f_j(x)$ using our abstract duals. The verification condition implemented in our tool is:

$$M(x_0) > \epsilon \cdot \sup_{x \in \mathbb{B}} \|\nabla m(x)\|_1 \quad (7)$$

In our implementation, we conservatively estimate this by summing the worst-case L1 norms of the gradients for the correct and competing classes:

$$\text{Certified Variation} = \epsilon \cdot \sum_i \max(|g_{low}^{(i)}|, |g_{high}^{(i)}|) \quad (8)$$

If the initial margin exceeds twice this variation (accounting for the worst-case drop in f_c and rise in f_j), the network is provably robust.

197 **Algorithm 1** Robustness Verification via Abstract Duals

198 **Require:** Neural Network f , Input x_0 , Label y , Radius ϵ
 199 **Ensure:** **True** if robust, **False** otherwise
 200 1: **Step 1: Abstract Initialization**
 201 2: $\hat{x}_{val} \leftarrow \text{AffineForm}(x_0 - \epsilon, x_0 + \epsilon)$ ▷ Input Box
 202 3: $\hat{x}_{grad} \leftarrow \text{Identity}(n_{in})$ ▷ Gradient w.r.t inputs
 203 4: $X \leftarrow \langle \hat{x}_{val}, \hat{x}_{grad} \rangle$
 204 5: **Step 2: Forward Propagation**
 205 6: **for** layer l in f **do**
 206 7: **if** l is $\text{Linear}(W, b)$ **then**
 207 8: $\hat{x}_{val} \leftarrow W\hat{x}_{val} + b$
 208 9: $\hat{x}_{grad} \leftarrow W\hat{x}_{grad}$
 209 10: **else if** l is $\text{Activation}(\sigma)$ **then**
 210 11: $c \leftarrow \text{center}(\hat{x}_{val})$
 211 12: $r \leftarrow \text{radius}(\hat{x}_{val})$
 212 13: $\epsilon_{err} \leftarrow 0.5 \cdot \max |\sigma''| \cdot r^2$ ▷ Curvature Error
 213 14: $\hat{x}_{val} \leftarrow \sigma(c) + \sigma'(c)(\hat{x}_{val} - c) + \epsilon_{err}$
 214 15: $\hat{\sigma}' \leftarrow \text{AffineApprox}(\sigma', \hat{x}_{val})$
 215 16: $\hat{x}_{grad} \leftarrow \hat{\sigma}' \otimes \hat{x}_{grad}$ ▷ Affine Multiplication
 216 17: **end if**
 217 18: **end for**
 218 19: **Step 3: Certification**
 219 20: $K \leftarrow \sup \|\hat{x}_{grad}\|_1$ ▷ Global Lipschitz Bound
 220 21: $M \leftarrow f(x_0)_y - \max_{j \neq y} f(x_0)_j$ ▷ Concrete Margin
 221 22: **if** $M > \epsilon \cdot K$ **then**
 222 23: **return True**
 223 24: **else**
 224 25: **return False**
 225 26: **end if**

229

3.3 Implementation and Stability Analysis

230 We give a functional implementation of our verification environment within OCaml, making use
 231 of a Monadic approach to cope with the computational graph. The distinct part of our tool is the
 232 *Gradient Stability Check*, used to identify the “wrapping effect” in real-time. The metric analyzes
 233 the intervals of the gradients on all dimensions corresponding to the inputs. If the interval $[g, \bar{g}]$
 234 strictly contains zero, it marks a neuron as unstable. This metric, $I_{unstable}$, highly correlates with
 235 failure cases of relational domains such as DeepPoly, providing a novel way of choosing verification
 236 heuristics.
 237

238

4 Formal Soundness of the Abstract Dual Domain

239 Here we establish the soundness of the Abstract Dual domain by demonstrating that the forward-
 240 mode propagation of abstract dual numbers yields a guaranteed over-approximation of both the
 241 output values and the Jacobian of a neural network for a specific input region. Consequently,
 242 the Lipschitz constant extracted from the final abstract gradient serves as a verified robustness
 243 certificate.
 244

246 4.1 Preliminaries

247 Let $f : \mathcal{R}^n \rightarrow \mathcal{R}^k$ be a feedforward neural network and let $X_0 \subseteq \mathcal{R}^n$ be a centrally symmetric input
 248 region. We utilize **Affine Arithmetic** for abstraction. An affine form \hat{z} represents a set of values
 249 parameterized by noise symbols $\epsilon_j \in [-1, 1]$:

$$250 \quad 251 \quad 252 \quad 253 \quad \hat{z} = \alpha_0 + \sum_{j=1}^m \alpha_j \epsilon_j \quad (9)$$

254 The **concretization** function $\gamma(\cdot)$ maps an affine form to its corresponding set of real values:

$$255 \quad 256 \quad 257 \quad 258 \quad \gamma(\hat{z}) \triangleq \left\{ \alpha_0 + \sum_{j=1}^m \alpha_j \epsilon_j \mid \forall j : \epsilon_j \in [-1, 1] \right\} \quad (10)$$

259 This definition extends element-wise to vectors and matrices.

260 We define the set of all concrete Jacobians reachable over the input region X_0 as:

$$261 \quad 262 \quad \mathcal{D}(f, X_0) \triangleq \{J_f(x) \in \mathcal{R}^{k \times n} \mid x \in X_0\} \quad (11)$$

263 *Definition 4.1 (Sound Abstract Dual).* An abstract dual number $\mathcal{X} \triangleq \langle \hat{x}_{val}, \hat{x}_{grad} \rangle$ is *sound* for a
 264 function f over X_0 if:

$$265 \quad 266 \quad \forall x \in X_0 : f(x) \in \gamma(\hat{x}_{val}), \quad (12)$$

$$267 \quad 268 \quad \forall x \in X_0 : J_f(x) \in \gamma(\hat{x}_{grad}) \quad (13)$$

269 4.2 Soundness of Linear Layers

270 *LEMMA 4.2 (LINEAR LAYER SOUNDNESS).* Let \mathcal{X} be a sound abstract dual for f over X_0 . Consider
 271 a linear layer $L(x) = Wx + b$. The abstract transformer $\mathcal{Y} \triangleq WX + b$ is sound for the composition
 272 $L \circ f$ over X_0 .

273 *PROOF.* The Jacobian of the affine transformation L is constant: $J_L(x) \equiv W$. By the multivariate
 274 chain rule, the Jacobian of the composition is:

$$275 \quad 276 \quad 277 \quad J_{L \circ f}(x) = J_L(f(x)) \cdot J_f(x) = W \cdot J_f(x) \quad (14)$$

278 The abstract transformer computes the output components as:

$$279 \quad 280 \quad \hat{y}_{val} = W\hat{x}_{val} + b, \quad (15)$$

$$281 \quad 282 \quad \hat{y}_{grad} = W\hat{x}_{grad} \quad (16)$$

283 **1. Value Soundness:** Since affine arithmetic is exact for linear operations, the concretization
 284 satisfies:

$$285 \quad 286 \quad \gamma(\hat{y}_{val}) = \{Wz + b \mid z \in \gamma(\hat{x}_{val})\} \quad (17)$$

287 Because \mathcal{X} is sound, $f(x) \in \gamma(\hat{x}_{val})$ for all $x \in X_0$. Therefore, $L(f(x)) \in \gamma(\hat{y}_{val})$.

288 **2. Gradient Soundness:** By the inductive hypothesis, $J_f(x) \in \gamma(\hat{x}_{grad})$ for all $x \in X_0$. Since
 289 matrix multiplication is a linear operation, the property of affine arithmetic ensures that:

$$290 \quad 291 \quad \forall J \in \gamma(\hat{x}_{grad}), \quad W \cdot J \in \gamma(W\hat{x}_{grad}) \quad (18)$$

292 Substituting the concrete Jacobian $J_f(x)$ for J , we obtain $W \cdot J_f(x) \in \gamma(\hat{y}_{grad})$. Thus, \mathcal{Y} is sound. \square

295 4.3 Soundness of Nonlinear Activations

296 LEMMA 4.3 (ACTIVATION FUNCTION SOUNDNESS). *Let $\sigma : \mathcal{R} \rightarrow \mathcal{R}$ be a differentiable activation
297 function applied element-wise. The abstract dual transformer, defined using interval bounds on the
298 derivative σ' , is sound.*

299 PROOF. Let $h(x) \triangleq \sigma(f(x))$. By the chain rule applied element-wise, the Jacobian is:

$$300 \quad J_h(x) = \text{diag}(\sigma'(f(x))) \cdot J_f(x) \quad (19)$$

302 Let the range of the value abstraction be $[l, u] = \text{Range}(\gamma(\hat{x}_{val}))$. We compute the interval
303 enclosure of the derivative over this range:

$$305 \quad \Sigma' \triangleq \left[\min_{z \in [l, u]} \sigma'(z), \max_{z \in [l, u]} \sigma'(z) \right] \quad (20)$$

307 **1. Value Soundness:** Standard abstract interpretation techniques for σ construct \hat{h}_{val} such that
308 $\sigma(\gamma(\hat{x}_{val})) \subseteq \gamma(\hat{h}_{val})$.

309 **2. Gradient Soundness:** Since \hat{x}_{val} is sound, $f(x) \in [l, u]$ for all $x \in X_0$. Consequently, the
310 concrete derivative $\sigma'(f(x))$ is contained within the interval Σ' . The abstract gradient is updated
311 via interval-affine multiplication:

$$312 \quad \hat{h}_{grad} \triangleq \Sigma' \otimes \hat{x}_{grad} \quad (21)$$

313 This operation is conservative: it accounts for every possible scaling factor in Σ' applied to every
314 affine form in \hat{x}_{grad} . Therefore:

$$316 \quad \text{diag}(\sigma'(f(x))) \cdot J_f(x) \in \gamma(\Sigma' \otimes \hat{x}_{grad}) \quad (22)$$

317 This confirms that $J_h(x) \in \gamma(\hat{h}_{grad})$, completing the proof. \square

319 4.4 Global Lipschitz Soundness

320 THEOREM 4.4 (SOUND LIPSCHITZ CERTIFICATION). *Let \hat{y}_{grad} be the abstract Jacobian at the network
321 output. The computed constant K_{comp} derived from \hat{y}_{grad} is a sound upper bound on the L_∞ -Lipschitz
322 constant of f over X_0 .*

323 PROOF. By induction on the network depth, the final abstract gradient satisfies:

$$325 \quad \forall x \in X_0 : \quad J_f(x) \in \gamma(\hat{y}_{grad}) \quad (23)$$

326 The local Lipschitz constant of f at x with respect to the L_∞ norm is the induced ∞ -norm of the
327 Jacobian:

$$328 \quad \|J_f(x)\|_\infty = \max_i \sum_j |(J_f(x))_{ij}| \quad (24)$$

330 The global Lipschitz constant is $K = \sup_{x \in X_0} \|J_f(x)\|_\infty$

331 Let the (i, j) -th entry of the abstract Jacobian matrix \hat{y}_{grad} be the affine form $\hat{g}_{ij} = c_0 + \sum_k c_k \epsilon_k$.
332 The maximum absolute value of this entry is bounded by its L_1 coefficient norm:

$$334 \quad \sup |\gamma(\hat{g}_{ij})| \leq |c_0| + \sum_k |c_k| \triangleq \mu_{ij} \quad (25)$$

336 We define the computed constant as the maximum row sum of these upper bounds $K_{comp} \triangleq$
337 $\max_i \sum_j \mu_{ij}$. By the triangle inequality, for any concrete $J \in \gamma(\hat{y}_{grad})$, $\|J\|_\infty \leq K_{comp}$. Since $J_f(x) \in$
338 $\gamma(\hat{y}_{grad})$, it follows that $K \leq K_{comp}$.

339 Finally, by the Mean Value Theorem, for any $x, x' \in X_0$:

$$340 \quad \|f(x) - f(x')\|_\infty \leq K_{comp} \cdot \|x - x'\|_\infty \quad (26)$$

342 Thus, K_{comp} is a valid Lipschitz certificate. \square

344 5 Evaluation

345 This section presents an evaluation of our method and contrasts its behavior with the DeepPoly
 346 abstraction. Concretely, we focus on the following research questions:

- 347 □ **RQ1 (Effectiveness):** In which settings does global gradient analysis yield tighter robustness
 348 certificates than layer-wise relational abstractions?
- 349 □ **RQ2 (Diagnostic Value):** Is the proposed notion of *gradient instability* predictive of verifi-
 350 cation failure?
- 351 □ **RQ3 (Scalability and Smoothness):** How does performance vary with network size and
 352 choice of activation function (Sigmoid versus ReLU)?

354 5.1 Experimental Setup

355 All experiments were implemented in **OCaml** (version 4.12+), building on the `ocaml-nn` library for
 356 neural network primitives. Evaluations were performed on a standard MacBook Pro equipped with
 357 an Apple M1 Pro processor.

359 *Benchmarks.* We consider fully connected classifiers trained on the MNIST dataset. To isolate
 360 the effect of network capacity, we evaluate three architectures of increasing width:

- 361 □ **Tiny Net:** 784 inputs → 2 hidden units → 10 outputs.
- 362 □ **Small Net:** 784 inputs → 5 hidden units → 10 outputs.
- 363 □ **Standard Net:** 784 inputs → 10 hidden units → 10 outputs.

365 All models were trained using RMSProp. Robustness was evaluated under L_∞ perturbations with
 366 radii $\epsilon \in \{0.01, 0.02, 0.03, 0.12\}$.

367 5.2 Results and Analysis

369 *5.2.1 RQ1: When Gradient Bounds Succeed and DeepPoly Fails.* A central empirical observation is
 370 the existence of a narrow regime in which the global gradient bound certifies robustness while
 371 DeepPoly does not. These cases arise almost exclusively for the **Tiny Net** architecture with Sigmoid
 372 activations at small perturbation radii.

373 Table 1. Cases in which the Gradient Method certifies robustness while DeepPoly fails.

376 Network	ϵ	Count	% of Total
Tiny Net	0.01	5	16.1%
Tiny Net	0.02	3	9.7%
Tiny Net	0.03	2	6.5%
Small Net	All	0	0.0%
Standard Net	All	0	0.0%

384 As shown in Table 1, at $\epsilon = 0.01$ the gradient-based method recovers over 16% of the test instances
 385 that DeepPoly fails to verify. These instances correspond to the following MNIST indices:

- 386 □ $\epsilon = 0.01$: indices 19, 76, 122, 139, 148 (all with true label 8);
- 387 □ $\epsilon = 0.02$: indices 122, 139, 148;
- 388 □ $\epsilon = 0.03$: indices 139, 148.

390 *Case Study: Index 139 (Label 8).* To better understand this discrepancy, we examine the instance
 391 at index 139 with $\epsilon = 0.03$.

- 393 **DeepPoly:** UNSTABLE. The forward-propagated margin collapses to an interval with effec-
394 tively zero width, indicating overlap with a competing class.
- 395 **Gradient Method:** ROBUST.
- 396 **Gradient Stability:** NO (480), indicating that 480 input dimensions exhibit sign changes
397 in the abstract gradient.

398 Despite substantial gradient instability, the *magnitude* of the gradient remains small. Conse-
399 quently, the certified variation $\epsilon \cdot \|\nabla f\|_1$ stays below the true margin between the predicted class
400 and its nearest competitor. In contrast, DeepPoly’s layer-wise abstraction appears to accumulate
401 sufficient over-approximation error—consistent with the classical wrapping effect—to lose the
402 certificate.

403 *Conclusion (RQ1).* Global gradient bounds can outperform DeepPoly in shallow, low-capacity
404 networks, where accumulated abstraction error dominates and local gradient magnitudes remain
405 small. This advantage disappears as depth or width increases: for both the Small and Standard
406 networks, DeepPoly strictly dominates, with no recovered cases.

407 **5.2.2 RQ2: Gradient Instability as a Predictor of Failure.** We next evaluate whether *gradient insta-*
408 *bility*—defined as the abstract gradient interval containing zero—serves as a reliable indicator of
409 verification difficulty. At the largest perturbation radius $\epsilon = 0.12$, we observe a **perfect correlation**:
410 every robustness failure coincides with gradient instability, across all architectures.

411 From a geometric perspective, instability reflects non-monotonicity of the decision function
412 within the input ball. In such regions, first-order linear approximations become inherently loose,
413 explaining the simultaneous failure of both DeepPoly and the Abstract Dual domain. These re-
414 sults support the use of the instability index I_{unstable} as a practical diagnostic tool for identifying
415 inherently hard instances.

416 **5.2.3 RQ3: Effect of Activation Smoothness.** Finally, we repeat the experiments using ReLU acti-
417 vations. In this setting, DeepPoly produces tighter or equal bounds in **all** cases. The degradation
418 of the Abstract Dual domain can be traced directly to the discontinuous derivative of ReLU. Any
419 unstable ReLU unit forces the abstract derivative to range over $[0, 1]$, dramatically inflating the
420 global Lipschitz constant. By contrast, Sigmoid activations admit smooth, bounded derivatives,
421 enabling significantly tighter gradient abstractions.

422 Overall, the evaluation shows that DeepPoly remains the most robust general-purpose abstrac-
423 tion. However, the Abstract Dual domain offers a complementary perspective: it is particularly
424 effective for *shallow networks with smooth activations*, where it can bypass the cumulative over-
425 approximation inherent to layer-wise methods. Moreover, gradient instability emerges as a mean-
426 ingful geometric signal, shedding light on when and why verification is likely to fail.

430 6 Related Work

431 The landscape of neural network verification is broadly divided into complete methods, which
432 provide exact certification at the cost of scalability, and incomplete methods, which utilize abstract
433 interpretation to achieve efficiency.

434 Initial efforts in the field prioritized exactness. Katz et al. [4] developed Reluplex, an SMT-based
435 solver that extends the Simplex algorithm to manage the piecewise linear nature of ReLU networks.
436 In a similar vein, Tjeng et al. [7] modeled the verification problem using Mixed Integer Linear
437 Programming (MIPVerify) to determine precise adversarial boundaries. Although these approaches
438 guarantee soundness and completeness, their NP-complete complexity renders them intractable
439 for modern, large-scale architectures.

To overcome the scalability bottlenecks of complete verifiers, Abstract Interpretation [1] has been widely adopted. Gehr et al. [2] introduced AI2, a framework employing zonotopes for bound propagation. This was advanced by Singh et al. [6] with DeepPoly, which combines intervals with polyhedral constraints to handle non-linearities like Sigmoid and Tanh through linear relaxation. While DeepPoly significantly improves scalability, it suffers from the "wrapping effect," where approximation errors accumulate with network depth. Weng et al. [8] similarly utilized linear bound propagation to expedite robustness certification.

While the theory of dual numbers is well-established for derivative evaluation [3], its utility in certifying robustness has been limited. Unlike purely geometric abstractions, our method uses Abstract Duals to compute global Lipschitz bounds. This allows us to bypass layer-wise error accumulation in specific regimes, offering a distinct advantage in shallow networks where gradient information remains coherent.

7 Conclusion

In this work, we explored the power of using Abstract Dual Numbers to certify the local robustness of neural networks. By lifting automatic differentiation into the abstract domain, we derived a method to compute global Lipschitz bounds that can certify invariance without explicit layer-by-layer geometric propagation.

Our comparative analysis with DeepPoly revealed a nuanced landscape. While DeepPoly remains the superior choice for general-purpose verification due to its ability to handle deep dependencies, our Gradient Method identified a specific "blind spot" in the relational approach. In shallow, smooth networks (Tiny Net with Sigmoid), our method successfully certified 16.1% of cases at $\epsilon = 0.01$ that DeepPoly failed to verify. This suggests that for low-depth architectures, the global gradient bound can be tighter than the accumulated over-approximation error of layer-wise abstract interpretation.

Finally, we established a strong correlation between *Gradient Instability* and robustness failure. In high-perturbation regimes, the presence of zero in the gradient interval served as a perfect predictor for the inability to certify robustness, highlighting its value as a diagnostic metric for detecting the breakdown of local linearity.

References

- [1] Patrick Cousot and Radhia Cousot. 1977. Abstract interpretation. *Symposium on Principles of Programming Languages* (Jan 1977). doi:10.1145/512950.512973
- [2] Timon Gehr, Matthew Mirman, Dana Drachsler-Cohen, Petar Tsankov, Swarat Chaudhuri, and Martin T. Vechev. 2018. AI2: Safety and Robustness Certification of Neural Networks with Abstract Interpretation. *2018 IEEE Symposium on Security and Privacy (SP)* (2018), 3–18. <https://api.semanticscholar.org/CorpusID:206579396>
- [3] Andreas Griewank and Andrea Walther. 2008. *Evaluating derivatives: principles and techniques of algorithmic differentiation*. SIAM.
- [4] Guy Katz, Clark Barrett, David L Dill, Kyle Julian, and Mykel J Kochenderfer. 2017. Reluplex: An efficient SMT solver for verifying deep neural networks. In *International Conference on Computer Aided Verification*. Springer, 97–117.
- [5] Gagandeep Singh, Timon Gehr, Matthew Mirman, Markus Püschel, and Martin Vechev. 2018. Fast and Effective Robustness Certification. In *Advances in Neural Information Processing Systems*, S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett (Eds.), Vol. 31. Curran Associates, Inc. https://proceedings.neurips.cc/paper_files/paper/2018/file/f2f446980d8e971ef3da97af089481c3-Paper.pdf
- [6] Gagandeep Singh, Timon Gehr, Markus Püschel, and Martin Vechev. 2019. An abstract domain for certifying neural networks, Vol. 3. 1–30. doi:10.1145/3290354
- [7] Vincent Tjeng, Kai Xiao, and Russ Tedrake. 2019. Evaluating Robustness of Neural Networks with Mixed Integer Programming. In *International Conference on Learning Representations*.
- [8] Tsui-Wei Weng, Huan Zhang, Hongge Chen, Zhao Song, Cho-Jui Hsieh, Duane S. Boning, Inderjit S. Dhillon, and Luca Daniel. 2018. Towards Fast Computation of Certified Robustness for ReLU Networks. *ArXiv* abs/1804.09699 (2018). <https://api.semanticscholar.org/CorpusID:13750928>