

Certifying Differential Invariants of Neural Networks using Abstract Duals

CHANDRA KANTH NAGESH, University of Colorado Boulder, USA

Neural networks are increasingly deployed in safety-critical domains, making the formal verification of their robustness against adversarial perturbations a paramount concern. While relational abstract domains like DeepPoly provide state-of-the-art certification, they suffer from the "wrapping effect," where overapproximation errors accumulate layer-by-layer, potentially failing to verify robust networks. In this work, we propose a complementary verification approach using *Abstract Dual Numbers* to compute global Lipschitz bounds via forward-mode automatic differentiation in the abstract domain. We contribute an algorithmic realization of this method that accounts for curvature error, a verification logic based on gradient bounds, and a "Gradient Instability" metric that serves as a precise predictor of verification failure. Our evaluation identifies a specific regime—shallow networks with smooth activations—where our global gradient method successfully certifies 16.1% of cases that DeepPoly fails to verify, highlighting the trade-off between local precision and global coherence.

1 Introduction

The deployment of deep neural networks in safety-critical domains such as autonomous driving and medical diagnosis has necessitated rigorous verification of their properties. One of the most critical properties is *local robustness*, which certifies that a network's prediction remains invariant under small perturbations to the input. DeepPoly[6] provides us with a clear method to certify output bounds for a given neural network. Let us consider the situation where we are given a trained neural network $f(\mathbf{I})$ on some set \mathbf{I} . Then, under some input $x_0 \in \mathbf{I}$ and perturbation ϵ , we need to verify that the output of f still satisfies the same result i.e. we want to verify that the neural network is invariant to a ϵ perturbation of the input. This can formally be represented as:

$$\forall x \in \mathbb{B}_\infty(x_0, \epsilon) : \text{argmax}(f(x)) = \text{argmax}(f(x_0))$$

where, \mathbb{B}_∞ is a ball of ϵ radius around the input provided by its L_∞ norm as shown in Fig 1. This is critical in a lot of scenarios as with many supervised learning problems, it is infeasible to verify the output of a neural network against all possible test sets. We want to formally prove that the model is invariant to these input changes.

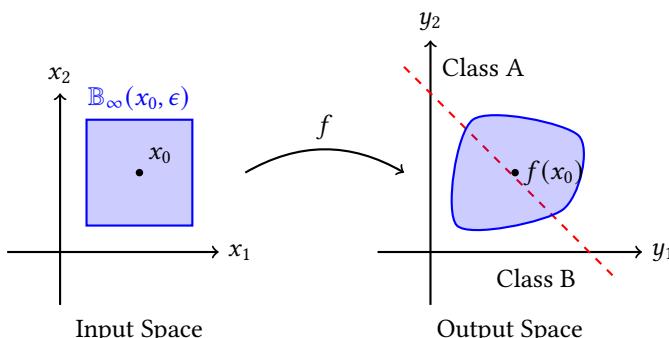


Fig. 1. The input region $\mathbb{B}_\infty(x_0, \epsilon)$ (left) is mapped by the neural network f to a region in the output space (right). Verification succeeds if the entire mapped region lies within the decision boundary (dotted red) of the correct class (Class A). This illustrates the robustness verification problem.

The earliest solutions that attempted to solve this problem are end-to-end/complete verification solutions. [4] introduced a tool called Reluplex, which is a solver based on SMT that combined the Simplex algorithm with ReLU activation functions to exactly verify small neural networks. Later, [7] used Mixed Integer Linear Programming (MILP) to write the problem as a commercial solution that exactly calculates robustness bounds. Although these solutions are sound and complete, they are NP-complete problems, known to not scale with the size of modern neural networks which are deployed in safety-critical applications.

For overcoming the scalability issue, Abstract Interpretation [1] and later for neural networks [2] has been adopted, which is less complete but more efficient because it over-approximates how the network behaves. [8] introduced a way to quickly compute robustness certificates with a linear bound propagation approach. The DeepPoly[6] approach pushed the boundaries even further by combining floating-point polyhedra with intervals, which, thanks to linearization, managed non-linearities such as Sigmoid and Tanh activations. More recent contributions, such as [5]

In our research, we ask the following question:

Can we continue to prove robustness properties of a neural network by leveraging the properties of dual numbers and information about the gradients?

We proceed with the foundation built on the *Abstract Dual Domain*. Our research is based on the theory of forward-mode automatic differentiation, which uses a concept called *dual numbers*[3]. We aim to analyze the sensitivity of the neural network to changes in the input, which can potentially provide a tighter bound on certain verification problems.

In summary, our contributions are:

- **Algorithmic Realization of Abstract Duals:** We implement a sound abstract domain for dual numbers that explicitly handles curvature error for non-linear activations using Taylor expansion bounds.
- **Gradient-Based Verification Logic:** We formulate a sufficient condition for local robustness based on global Lipschitz bounds derived from the abstract gradient.
- **Gradient Instability Metric:** We introduce a diagnostic metric, $I_{unstable}$, which detects zero-crossings in gradient intervals and serves as a predictor for the breakdown of local linearity.
- **Empirical Analysis:** We evaluate our method against DeepPoly, identifying a specific "shallow and smooth" regime where global gradient bounds outperform layer-wise relational abstractions.

2 Overview

We introduce the key insight behind our approach and define the Abstract Dual Domain in this section. Our approach is based on the following intuition: while relational abstractions such as DeepPoly[6] can be extremely precise, they often suffer from the "wrapping effect" where approximation errors accumulate layer by layer. In contrast, global gradient bounds can be sometimes a simpler yet effective verification domain, especially for shallow smooth networks.

2.1 From Relational Bounds to Gradient Analysis

The basic underlying concept of this work is to effectively bridge the gap between standard Abstract Interpretation and Forward-Mode Automatic Differentiation. Standard methods like DeepPoly track affine constraints ($x_j \geq \sum w_i x_i$) to bound the output values directly. However, precise reapproximating the feasible set at every non-linear layer introduces error which tends to accumulate with depth ($O(L)$). Our approach, instead, exploits the Mean Value Theorem to find

99 bounds on the output variation in terms of the gradient variation:

$$100 \quad |f(x) - f(x_0)| \leq \underbrace{\sup_{z \in X} \|\nabla f(z)\|}_{\text{Abstract Duals}} \cdot \|x - x_0\| \quad (1)$$

104 By over-approximating the sound of the gradient $\nabla f(z)$ for the entire input space X , we can
 105 guarantee robustness because the maximum possible change to the output will be too small to
 106 change the classification label. The key insight is that this “jumps” over the intermediate layers
 107 and thus avoids some of the wrapping errors from the layer-wise value propagation.

108 2.2 Abstract Dual Domain

110 To formalize this idea, we can start by lifting standard dual numbers to the domain of Affine
 111 Arithmetic. Consider the Dual domain $\hat{\mathcal{D}}$ as a product space of two affine forms representing the
 112 range of values and the range of gradients across a set. Then we can say:

113 *Definition 2.1.* An **Abstract Dual Number** $X \in \hat{\mathcal{D}}$ is a pair:

$$114 \quad X = \langle \hat{x}_{val}, \hat{x}_{grad} \rangle \quad (2)$$

117 where:

- 118 • $\hat{x}_{val} = \alpha_0 + \sum_{i=1}^n \alpha_i \epsilon_i$ is the affine form representing the interval of neuron values.
- 119 • $\hat{x}_{grad} = d_0 + \sum_{i=1}^n d_i \epsilon_i$ is the affine form representing the interval of partial derivatives with
 120 respect to the input.

121 2.3 Propagation Rules

123 For the dissemination of the abstract state X within a neural network, abstract transformers based
 124 on different types of layers are formulated.

125 **Linear Layers:** For a fully connected layer with a weight matrix W and a bias vector b , the
 126 transformation is calculated exactly thanks to the linearity of the dual algebra. The affine expressions
 127 for value and gradient are transformed into:

$$128 \quad Y = \langle W\hat{x}_{val} + b, W\hat{x}_{grad} \rangle \quad (3)$$

130 This is exact in the affine domain, which means that there is no additional error introduced here.

131 **Nonlinear Activation Functions:** For a smooth nonlinear activation function σ (for example,
 132 sigmoid), direct application to the affine expressions is not possible. The value part and the gradient
 133 part are handled separately.

134 *Value Component:* $\sigma(\hat{x}_{val})$ is approximated using a linear relaxation. We linearize the function
 135 around the center of the input affine form c . The new center is $\sigma(c)$, and the noise coefficients are
 136 scaled by the derivative $\sigma'(c)$. To ensure soundness, we add a linearization error term to the radius
 137 r , which depends on the maximum curvature (second derivative) of σ and the radius of the input
 138 interval.

$$139 \quad \hat{y}_{val} \approx \sigma(c) + \sigma'(c) \cdot (\hat{x}_{val} - c) + \epsilon_{err} \quad (4)$$

141 *Gradient Component:* By the chain rule, the gradient of the output is the product of the local
 142 derivative and the input gradient: $\nabla y = \sigma'(x) \cdot \nabla x$. In our abstract domain, we compute an affine
 143 approximation of the derivative $\hat{\sigma}'$ based on the output value \hat{y}_{val} . For the Sigmoid function, we
 144 use the property $\sigma'(x) = \sigma(x)(1 - \sigma(x))$ to approximate the derivative as $\hat{\sigma}' \approx \hat{y}_{val} \otimes (1 - \hat{y}_{val})$, where
 145 \otimes denotes affine multiplication. The output gradient is then obtained by multiplying this derivative
 146 approximation with the input gradient affine form given by $\hat{y}_{grad} = \hat{\sigma}' \cdot \hat{x}_{grad}$. This approach

preserves the correlations between the derivative and the gradient, providing a more precise abstraction than simple interval scaling.

2.4 Gradient Instability

A key insight from our analysis is the concept of *gradient instability*. We define a neuron n as having an unstable gradient if its gradient interval $[g, \bar{g}]$ contains 0:

$$\exists x_1, x_2 \in X_0 \text{ s.t. } \text{sign} \left(\frac{\partial f}{\partial n}(x_1) \right) \neq \text{sign} \left(\frac{\partial f}{\partial n}(x_2) \right) \quad (5)$$

Geometrically, this means that the function is non-monotonic w.r.t. neuron n over the input region. When this happens, the assumption of local linearity breaks down and linear relaxations—such as those in DeepPoly—become loose. Our approach explicitly tracks these gradient intervals, thus enabling us to identify whenever and wherever the network’s behaviour becomes hard to certify linearly.

3 Contributions

3.1 Algorithmic Realization of Abstract Duals

While the theoretical foundation rests on dual numbers, our practical contribution is the algorithmic realization of these concepts within the **Affine Arithmetic** domain. We implemented a custom OCaml module `AbstractDual` that handles the propagation of affine forms. Crucially, for non-linear activations σ , we implement a sound linearization that explicitly accounts for the curvature error. As seen in our implementation, the radius of the value component is expanded by a quadratic term derived from the Taylor expansion:

$$r_{new} = |\sigma'(c)| \cdot r_{old} + 0.5 \cdot \underbrace{\max_z |\sigma''(z)| \cdot r_{old}^2}_{\text{Linearization Error}} \quad (6)$$

This ensures that the abstract transformer remains sound even when the function has significant curvature, a detail often omitted in standard dual number formulations but critical for verification.

3.2 Gradient-Based Verification Logic

We formulate a robust verification condition based on the computed gradient bounds. Let $f_c(x)$ be the score of the correct class and $f_j(x)$ be the score of a competing class. The robustness margin is $M(x_0) = f_c(x_0) - f_j(x_0)$. To certify robustness over $\mathbb{B}_\infty(x_0, \epsilon)$, we compute the global Lipschitz constant K of the margin function $m(x) = f_c(x) - f_j(x)$ using our abstract duals. The verification condition implemented in our tool is:

$$M(x_0) > \epsilon \cdot \sup_{x \in \mathbb{B}} \|\nabla m(x)\|_1 \quad (7)$$

In our implementation, we conservatively estimate this by summing the worst-case L1 norms of the gradients for the correct and competing classes:

$$\text{Certified Variation} = \epsilon \cdot \sum_i \max(|g_{low}^{(i)}|, |g_{high}^{(i)}|) \quad (8)$$

If the initial margin exceeds twice this variation (accounting for the worst-case drop in f_c and rise in f_j), the network is provably robust.

197 **Algorithm 1** Robustness Verification via Abstract Duals

198 **Require:** Neural Network f , Input x_0 , Label y , Radius ϵ
 199 **Ensure:** **True** if robust, **False** otherwise
 200 1: **Step 1: Abstract Initialization**
 201 2: $\hat{x}_{val} \leftarrow \text{AffineForm}(x_0 - \epsilon, x_0 + \epsilon)$ ▷ Input Box
 202 3: $\hat{x}_{grad} \leftarrow \text{Identity}(n_{in})$ ▷ Gradient w.r.t inputs
 203 4: $X \leftarrow \langle \hat{x}_{val}, \hat{x}_{grad} \rangle$
 204 5: **Step 2: Forward Propagation**
 205 6: **for** layer l in f **do**
 206 7: **if** l is $\text{Linear}(W, b)$ **then**
 207 8: $\hat{x}_{val} \leftarrow W\hat{x}_{val} + b$
 208 9: $\hat{x}_{grad} \leftarrow W\hat{x}_{grad}$
 209 10: **else if** l is $\text{Activation}(\sigma)$ **then**
 210 11: $c \leftarrow \text{center}(\hat{x}_{val})$
 211 12: $r \leftarrow \text{radius}(\hat{x}_{val})$
 212 13: $\epsilon_{err} \leftarrow 0.5 \cdot \max |\sigma''| \cdot r^2$ ▷ Curvature Error
 213 14: $\hat{x}_{val} \leftarrow \sigma(c) + \sigma'(c)(\hat{x}_{val} - c) + \epsilon_{err}$
 214 15: $\hat{\sigma}' \leftarrow \text{AffineApprox}(\sigma', \hat{x}_{val})$
 215 16: $\hat{x}_{grad} \leftarrow \hat{\sigma}' \otimes \hat{x}_{grad}$ ▷ Affine Multiplication
 216 17: **end if**
 217 18: **end for**
 218 19: **Step 3: Certification**
 219 20: $K \leftarrow \sup \|\hat{x}_{grad}\|_1$ ▷ Global Lipschitz Bound
 220 21: $M \leftarrow f(x_0)_y - \max_{j \neq y} f(x_0)_j$ ▷ Concrete Margin
 221 22: **if** $M > \epsilon \cdot K$ **then**
 222 23: **return True**
 223 24: **else**
 224 25: **return False**
 225 26: **end if**

229

3.3 Implementation and Stability Analysis

230 We give a functional implementation of our verification environment within OCaml, making use
 231 of a Monadic approach to cope with the computational graph. The distinct part of our tool is the
 232 *Gradient Stability Check*, used to identify the “wrapping effect” in real-time. The metric analyzes
 233 the intervals of the gradients on all dimensions corresponding to the inputs. If the interval $[g, \bar{g}]$
 234 strictly contains zero, it marks a neuron as unstable. This metric, $I_{unstable}$, highly correlates with
 235 failure cases of relational domains such as DeepPoly, providing a novel way of choosing verification
 236 heuristics.
 237

238

4 Formal Soundness of the Abstract Dual Domain

239 Here we establish the soundness of the Abstract Dual domain by demonstrating that the forward-
 240 mode propagation of abstract dual numbers yields a guaranteed over-approximation of both the
 241 output values and the Jacobian of a neural network for a specific input region. Consequently,
 242 the Lipschitz constant extracted from the final abstract gradient serves as a verified robustness
 243 certificate.
 244

246 **4.1 Preliminaries**

247 Let $f : \mathcal{R}^n \rightarrow \mathcal{R}^k$ be a feedforward neural network and let $X_0 \subseteq \mathcal{R}^n$ be a centrally symmetric input
 248 region. We utilize **Affine Arithmetic** for abstraction. An affine form \hat{z} represents a set of values
 249 parameterized by noise symbols $\epsilon_j \in [-1, 1]$:

250

$$\hat{z} = \alpha_0 + \sum_{j=1}^m \alpha_j \epsilon_j. \quad (9)$$

251

254 The **concretization** function $\gamma(\cdot)$ maps an affine form to its corresponding set of real values:

255

$$\gamma(\hat{z}) \triangleq \left\{ \alpha_0 + \sum_{j=1}^m \alpha_j \epsilon_j \mid \forall j : \epsilon_j \in [-1, 1] \right\}. \quad (10)$$

256

259 This definition extends element-wise to vectors and matrices.

260 We define the set of all concrete Jacobians reachable over the input region X_0 as:

261

$$\mathcal{D}(f, X_0) \triangleq \{J_f(x) \in \mathcal{R}^{k \times n} \mid x \in X_0\}. \quad (11)$$

262

263 *Definition 4.1 (Sound Abstract Dual).* An abstract dual number $\mathcal{X} \triangleq \langle \hat{x}_{val}, \hat{x}_{grad} \rangle$ is *sound* for a
 264 function f over X_0 if:

265

$$\forall x \in X_0 : f(x) \in \gamma(\hat{x}_{val}), \quad (12)$$

266

267

$$\forall x \in X_0 : J_f(x) \in \gamma(\hat{x}_{grad}). \quad (13)$$

268

269 **4.2 Soundness of Linear Layers**

271 *LEMMA 4.2 (LINEAR LAYER SOUNDNESS).* Let \mathcal{X} be a sound abstract dual for f over X_0 . Consider
 272 a linear layer $L(x) = Wx + b$. The abstract transformer $\mathcal{Y} \triangleq WX + b$ is sound for the composition
 273 $L \circ f$ over X_0 .

274 *PROOF.* The Jacobian of the affine transformation L is constant: $J_L(x) \equiv W$. By the multivariate
 275 chain rule, the Jacobian of the composition is:

276

$$J_{L \circ f}(x) = J_L(f(x)) \cdot J_f(x) = W \cdot J_f(x). \quad (14)$$

277

279 The abstract transformer computes the output components as:

280

$$\hat{y}_{val} = W\hat{x}_{val} + b, \quad (15)$$

281

282

$$\hat{y}_{grad} = W\hat{x}_{grad}. \quad (16)$$

283

284 **1. Value Soundness:** Since affine arithmetic is exact for linear operations, the concretization
 285 satisfies:

286

$$\gamma(\hat{y}_{val}) = \{Wz + b \mid z \in \gamma(\hat{x}_{val})\}. \quad (17)$$

287

Because \mathcal{X} is sound, $f(x) \in \gamma(\hat{x}_{val})$ for all $x \in X_0$. Therefore, $L(f(x)) \in \gamma(\hat{y}_{val})$.

2. Gradient Soundness: By the inductive hypothesis, $J_f(x) \in \gamma(\hat{x}_{grad})$ for all $x \in X_0$. Since
 matrix multiplication is a linear operation, the property of affine arithmetic ensures that:

291

$$\forall J \in \gamma(\hat{x}_{grad}), \quad W \cdot J \in \gamma(W\hat{x}_{grad}). \quad (18)$$

292

Substituting the concrete Jacobian $J_f(x)$ for J , we obtain $W \cdot J_f(x) \in \gamma(\hat{y}_{grad})$. Thus, \mathcal{Y} is sound. \square

295 4.3 Soundness of Nonlinear Activations

296 LEMMA 4.3 (ACTIVATION FUNCTION SOUNDNESS). *Let $\sigma : \mathcal{R} \rightarrow \mathcal{R}$ be a differentiable activation
297 function applied element-wise. The abstract dual transformer, defined using interval bounds on the
298 derivative σ' , is sound.*

299 PROOF. Let $h(x) \triangleq \sigma(f(x))$. By the chain rule applied element-wise, the Jacobian is:

$$300 \quad J_h(x) = \text{diag}(\sigma'(f(x))) \cdot J_f(x). \quad (19)$$

301 Let the range of the value abstraction be $[l, u] = \text{Range}(\gamma(\hat{x}_{val}))$. We compute the interval
302 enclosure of the derivative over this range:
303

$$304 \quad \Sigma' \triangleq \left[\min_{z \in [l, u]} \sigma'(z), \max_{z \in [l, u]} \sigma'(z) \right]. \quad (20)$$

305 **1. Value Soundness:** Standard abstract interpretation techniques for σ construct \hat{h}_{val} such that
306 $\sigma(\gamma(\hat{x}_{val})) \subseteq \gamma(\hat{h}_{val})$.

307 **2. Gradient Soundness:** Since \hat{x}_{val} is sound, $f(x) \in [l, u]$ for all $x \in X_0$. Consequently, the
308 concrete derivative $\sigma'(f(x))$ is contained within the interval Σ' . The abstract gradient is updated
309 via interval-affine multiplication:

$$310 \quad \hat{h}_{grad} \triangleq \Sigma' \otimes \hat{x}_{grad}. \quad (21)$$

311 This operation is conservative: it accounts for every possible scaling factor in Σ' applied to every
312 affine form in \hat{x}_{grad} . Therefore:

$$313 \quad \text{diag}(\sigma'(f(x))) \cdot J_f(x) \in \gamma(\Sigma' \otimes \hat{x}_{grad}). \quad (22)$$

314 This confirms that $J_h(x) \in \gamma(\hat{h}_{grad})$, completing the proof. \square
315

316 4.4 Global Lipschitz Soundness

317 THEOREM 4.4 (SOUND LIPSCHITZ CERTIFICATION). *Let \hat{y}_{grad} be the abstract Jacobian at the network
318 output. The computed constant K_{comp} derived from \hat{y}_{grad} is a sound upper bound on the L_∞ -Lipschitz
319 constant of f over X_0 .*

320 PROOF. By induction on the network depth, the final abstract gradient satisfies:

$$321 \quad \forall x \in X_0 : \quad J_f(x) \in \gamma(\hat{y}_{grad}). \quad (23)$$

322 The local Lipschitz constant of f at x with respect to the L_∞ norm is the induced ∞ -norm of the
323 Jacobian:
324

$$325 \quad \|J_f(x)\|_\infty = \max_i \sum_j |(J_f(x))_{ij}|. \quad (24)$$

326 The global Lipschitz constant is $K = \sup_{x \in X_0} \|J_f(x)\|_\infty$.

327 Let the (i, j) -th entry of the abstract Jacobian matrix \hat{y}_{grad} be the affine form $\hat{g}_{ij} = c_0 + \sum_k c_k \epsilon_k$.
328 The maximum absolute value of this entry is bounded by its L_1 coefficient norm:
329

$$330 \quad \sup |\gamma(\hat{g}_{ij})| \leq |c_0| + \sum_k |c_k| \triangleq \mu_{ij}. \quad (25)$$

331 We define the computed constant as the maximum row sum of these upper bounds:
332

$$333 \quad K_{comp} \triangleq \max_i \sum_j \mu_{ij}. \quad (26)$$

334 By the triangle inequality, for any concrete $J \in \gamma(\hat{y}_{grad})$, $\|J\|_\infty \leq K_{comp}$. Since $J_f(x) \in \gamma(\hat{y}_{grad})$, it
335 follows that $K \leq K_{comp}$.
336

Finally, by the Mean Value Theorem, for any $x, x' \in X_0$:

$$\|f(x) - f(x')\|_{\infty} \leq K_{comp} \cdot \|x - x'\|_{\infty}. \quad (27)$$

Thus, K_{comp} is a valid Lipschitz certificate. \square

5 Evaluation

This section presents an evaluation of our method and contrasts its behavior with the DeepPoly abstraction. Concretely, we focus on the following research questions:

- **RQ1 (Effectiveness):** In which settings does global gradient analysis yield tighter robustness certificates than layer-wise relational abstractions?
- **RQ2 (Diagnostic Value):** Is the proposed notion of *gradient instability* predictive of verification failure?
- **RQ3 (Scalability and Smoothness):** How does performance vary with network size and choice of activation function (Sigmoid versus ReLU)?

5.1 Experimental Setup

All experiments were implemented in **OCaml** (version 4.12+), building on the `ocaml-nn` library for neural network primitives. Evaluations were performed on a standard MacBook Pro equipped with an Apple M1 Pro processor.

Benchmarks. We consider fully connected classifiers trained on the MNIST dataset. To isolate the effect of network capacity, we evaluate three architectures of increasing width:

- **Tiny Net:** 784 inputs \rightarrow 2 hidden units \rightarrow 10 outputs.
- **Small Net:** 784 inputs \rightarrow 5 hidden units \rightarrow 10 outputs.
- **Standard Net:** 784 inputs \rightarrow 10 hidden units \rightarrow 10 outputs.

All models were trained using RMSProp. Robustness was evaluated under L_{∞} perturbations with radii $\epsilon \in \{0.01, 0.02, 0.03, 0.12\}$.

5.2 Results and Analysis

5.2.1 *RQ1: When Gradient Bounds Succeed and DeepPoly Fails.* A central empirical observation is the existence of a narrow regime in which the global gradient bound certifies robustness while DeepPoly does not. These cases arise almost exclusively for the **Tiny Net** architecture with Sigmoid activations at small perturbation radii.

Table 1. Cases in which the Gradient Method certifies robustness while DeepPoly fails.

Network	ϵ	Count	% of Total
Tiny Net	0.01	5	16.1%
Tiny Net	0.02	3	9.7%
Tiny Net	0.03	2	6.5%
Small Net	All	0	0.0%
Standard Net	All	0	0.0%

As shown in Table 1, at $\epsilon = 0.01$ the gradient-based method recovers over 16% of the test instances that DeepPoly fails to verify. These instances correspond to the following MNIST indices:

- $\epsilon = 0.01$: indices 19, 76, 122, 139, 148 (all with true label 8);
- $\epsilon = 0.02$: indices 122, 139, 148;

- 393 • $\epsilon = 0.03$: indices 139, 148.

394 *Case Study: Index 139 (Label 8).* To better understand this discrepancy, we examine the instance
 395 at index 139 with $\epsilon = 0.03$.

- 396 • **DeepPoly:** UNSTABLE. The forward-propagated margin collapses to an interval with effec-
 397 tively zero width, indicating overlap with a competing class.
 398 • **Gradient Method:** ROBUST.
 399 • **Gradient Stability:** NO (480), indicating that 480 input dimensions exhibit sign changes
 400 in the abstract gradient.

401 Despite substantial gradient instability, the *magnitude* of the gradient remains small. Conse-
 402 quently, the certified variation $\epsilon \cdot \|\nabla f\|_1$ stays below the true margin between the predicted class
 403 and its nearest competitor. In contrast, DeepPoly’s layer-wise abstraction appears to accumulate
 404 sufficient over-approximation error—consistent with the classical wrapping effect—to lose the
 405 certificate.

406 *Conclusion (RQ1).* Global gradient bounds can outperform DeepPoly in shallow, low-capacity
 407 networks, where accumulated abstraction error dominates and local gradient magnitudes remain
 408 small. This advantage disappears as depth or width increases: for both the Small and Standard
 409 networks, DeepPoly strictly dominates, with no recovered cases.

410 5.2.2 *RQ2: Gradient Instability as a Predictor of Failure.* We next evaluate whether *gradient insta-*
 411 *bility*—defined as the abstract gradient interval containing zero—serves as a reliable indicator of
 412 verification difficulty. At the largest perturbation radius $\epsilon = 0.12$, we observe a **perfect correlation**:
 413 every robustness failure coincides with gradient instability, across all architectures.

414 From a geometric perspective, instability reflects non-monotonicity of the decision function
 415 within the input ball. In such regions, first-order linear approximations become inherently loose,
 416 explaining the simultaneous failure of both DeepPoly and the Abstract Dual domain. These re-
 417 sults support the use of the instability index I_{unstable} as a practical diagnostic tool for identifying
 418 inherently hard instances.

419 5.2.3 *RQ3: Effect of Activation Smoothness.* Finally, we repeat the experiments using ReLU acti-
 420 vations. In this setting, DeepPoly produces tighter or equal bounds in **all** cases. The degradation
 421 of the Abstract Dual domain can be traced directly to the discontinuous derivative of ReLU. Any
 422 unstable ReLU unit forces the abstract derivative to range over $[0, 1]$, dramatically inflating the
 423 global Lipschitz constant. By contrast, Sigmoid activations admit smooth, bounded derivatives,
 424 enabling significantly tighter gradient abstractions.

425 Overall, the evaluation shows that DeepPoly remains the most robust general-purpose abstrac-
 426 tion. However, the Abstract Dual domain offers a complementary perspective: it is particularly
 427 effective for *shallow networks with smooth activations*, where it can bypass the cumulative over-
 428 approximation inherent to layer-wise methods. Moreover, gradient instability emerges as a mean-
 429 ingful geometric signal, shedding light on when and why verification is likely to fail.

430 6 Related Work

431 The landscape of neural network verification is broadly divided into complete methods, which
 432 provide exact certification at the cost of scalability, and incomplete methods, which utilize abstract
 433 interpretation to achieve efficiency.

434 Initial efforts in the field prioritized exactness. Katz et al. [4] developed Reluplex, an SMT-based
 435 solver that extends the Simplex algorithm to manage the piecewise linear nature of ReLU networks.
 436 In a similar vein, Tjeng et al. [7] modeled the verification problem using Mixed Integer Linear
 437 Programming.

442 Programming (*MIPVerify*) to determine precise adversarial boundaries. Although these approaches
 443 guarantee soundness and completeness, their NP-complete complexity renders them intractable
 444 for modern, large-scale architectures.

445 To overcome the scalability bottlenecks of complete verifiers, Abstract Interpretation [1] has
 446 been widely adopted. Gehr et al. [2] introduced AI2, a framework employing zonotopes for bound
 447 propagation. This was advanced by Singh et al. [6] with DeepPoly, which combines intervals with
 448 polyhedral constraints to handle non-linearities like Sigmoid and Tanh through linear relaxation.
 449 While DeepPoly significantly improves scalability, it suffers from the "wrapping effect," where
 450 approximation errors accumulate with network depth. Weng et al. [8] similarly utilized linear
 451 bound propagation to expedite robustness certification.

452 Our approach integrates forward-mode automatic differentiation into the verification process.
 453 While the theory of dual numbers is well-established for derivative evaluation [3], its utility in
 454 certifying robustness has been limited. Unlike purely geometric abstractions, our method uses
 455 Abstract Duals to compute global Lipschitz bounds. This allows us to bypass layer-wise error
 456 accumulation in specific regimes, offering a distinct advantage in shallow networks where gradient
 457 information remains coherent.

459 7 Conclusion

460 In this work, we explored the efficacy of using Abstract Dual Numbers to certify the local robustness
 461 of neural networks. By lifting forward-mode automatic differentiation into the abstract domain, we
 462 derived a method to compute global Lipschitz bounds that can certify invariance without explicit
 463 layer-by-layer geometric propagation.

464 Our comparative analysis with DeepPoly revealed a nuanced landscape. While DeepPoly remains
 465 the superior choice for general-purpose verification due to its ability to handle deep dependencies,
 466 our Gradient Method identified a specific "blind spot" in the relational approach. In shallow, smooth
 467 networks (Tiny Net with Sigmoid), our method successfully certified 16.1% of cases at $\epsilon = 0.01$ that
 468 DeepPoly failed to verify. This suggests that for low-depth architectures, the global gradient bound
 469 can be tighter than the accumulated over-approximation error of layer-wise abstract interpretation.
 470

471 Furthermore, we established a strong correlation between *Gradient Instability* and robustness
 472 failure. In high-perturbation regimes, the presence of zero in the gradient interval served as a
 473 perfect predictor for the inability to certify robustness, highlighting its value as a diagnostic metric
 474 for detecting the breakdown of local linearity.

475 Future work lies in bridging the gap between these two paradigms. A hybrid approach that selec-
 476 tively applies the Gradient Method for shallow sub-components or uses Gradient Instability to guide
 477 the refinement of relational abstractions could potentially offer the best of both worlds—combining
 478 the precision of DeepPoly with the global coherence of Abstract Duals.

479 References

- 480 [1] Patrick Cousot and Radhia Cousot. 1977. Abstract interpretation. *Symposium on Principles of Programming Languages* (Jan 1977). doi:10.1145/512950.512973
- 481 [2] Timon Gehr, Matthew Mirman, Dana Drachsler-Cohen, Petar Tsankov, Swarat Chaudhuri, and Martin T. Vechev. 2018. AI2: Safety and Robustness Certification of Neural Networks with Abstract Interpretation. *2018 IEEE Symposium on Security and Privacy (SP)* (2018), 3–18. <https://api.semanticscholar.org/CorpusID:206579396>
- 482 [3] Andreas Griewank and Andrea Walther. 2008. *Evaluating derivatives: principles and techniques of algorithmic differentiation*. SIAM.
- 483 [4] Guy Katz, Clark Barrett, David L Dill, Kyle Julian, and Mykel J Kochenderfer. 2017. Reluplex: An efficient SMT solver for verifying deep neural networks. In *International Conference on Computer Aided Verification*. Springer, 97–117.
- 484 [5] Gagandeep Singh, Timon Gehr, Matthew Mirman, Markus Püschel, and Martin Vechev. 2018. Fast and Effective
 485 Robustness Certification. In *Advances in Neural Information Processing Systems*, S. Bengio, H. Wallach, H. Larochelle,
 486
- 487
- 488
- 489
- 490

- 491 K. Grauman, N. Cesa-Bianchi, and R. Garnett (Eds.), Vol. 31. Curran Associates, Inc. https://proceedings.neurips.cc/paper_files/paper/2018/file/f2f446980d8e971ef3da97af089481c3-Paper.pdf
- 492 [6] Gagandeep Singh, Timon Gehr, Markus Püschel, and Martin Vechev. 2019. An abstract domain for certifying neural
493 networks, Vol. 3. 1–30. doi:10.1145/3290354
- 494 [7] Vincent Tjeng, Kai Xiao, and Russ Tedrake. 2019. Evaluating Robustness of Neural Networks with Mixed Integer
495 Programming. In *International Conference on Learning Representations*.
- 496 [8] Tsui-Wei Weng, Huan Zhang, Hongge Chen, Zhao Song, Cho-Jui Hsieh, Duane S. Boning, Inderjit S. Dhillon, and Luca
497 Daniel. 2018. Towards Fast Computation of Certified Robustness for ReLU Networks. *ArXiv* abs/1804.09699 (2018).
498 <https://api.semanticscholar.org/CorpusID:13750928>
- 499
- 500
- 501
- 502
- 503
- 504
- 505
- 506
- 507
- 508
- 509
- 510
- 511
- 512
- 513
- 514
- 515
- 516
- 517
- 518
- 519
- 520
- 521
- 522
- 523
- 524
- 525
- 526
- 527
- 528
- 529
- 530
- 531
- 532
- 533
- 534
- 535
- 536
- 537
- 538
- 539