

# Certifying Differential Invariants of Neural Networks using Abstract Duals

CHANDRA KANTH NAGESH, University of Colorado Boulder, USA

Neural networks are increasingly used in safety-critical domains necessitating formal guarantees about their properties under perturbations to inputs. Existing robust verification techniques, typified by DeepPoly primarily focus on robustness analysis during forward mode where certifying output stability is given by  $f(\mathbf{I}) \subseteq [y_L, y_R]$  for an  $L_\infty$ -constrained input set  $\mathbf{I} = \{\mathbf{x} \mid \|\mathbf{x} - \mathbf{x}_0\|_\infty \leq \epsilon\}$ . While such methods employ a sophisticated polyhedral abstract domain (combining intervals and affine forms) to generate sound, tight bounds, this entire class of analysis still fails to provide sound guarantees over the behaviour. This oversight creates a critical verification gap for gradient-dependent systems where the Backward Pass Verification Problem which can be described formally as bounding the Jacobian  $J(\mathbf{x}) = \nabla_{\mathbf{x}} f(\mathbf{x})$  such that  $J(\mathbf{x}) \subseteq [J_L, J_R]$  for all  $\mathbf{x} \in \mathbf{I}$  is essential.

## 1 Introduction

The deployment of deep neural networks in safety-critical domains such as autonomous driving and medical diagnosis has necessitated rigorous verification of their properties. One of the most critical properties is *local robustness*, which certifies that a network's prediction remains invariant under small perturbations to the input. DeepPoly[6] provides us with a clear method to certify output bounds for a given neural network. Let us consider the situation where we are given a trained neural network  $f(\mathbf{I})$  on some set  $\mathbf{I}$ . Then, under some input  $x_0 \in \mathbf{I}$  and perturbation  $\epsilon$ , we need to verify that the output of  $f$  still satisfies the same result i.e. we want to verify that the neural network is invariant to a  $\epsilon$  perturbation of the input. This can formally be represented as:

$$\forall x \in \mathbb{B}_\infty(x_0, \epsilon) : \text{argmax}(f(x)) = \text{argmax}(f(x_0))$$

where,  $\mathbb{B}_\infty$  is a ball of  $\epsilon$  radius around the input provided by its  $L_\infty$  norm as shown in Fig 1. This is critical in a lot of scenarios as with many supervised learning problems, it is infeasible to verify the output of a neural network against all possible test sets. We want to formally prove that the model is invariant to these input changes.

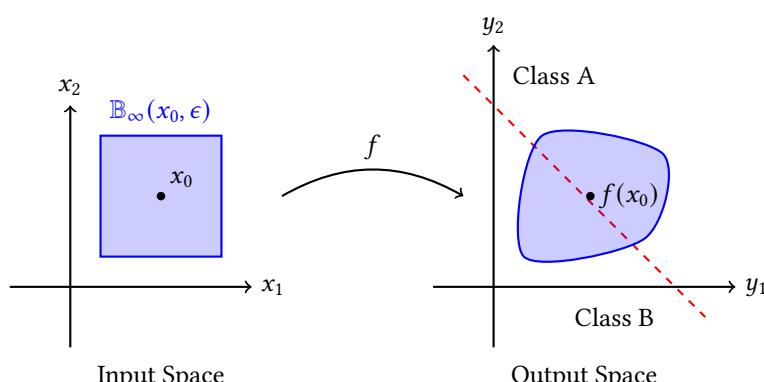


Fig. 1. The input region  $\mathbb{B}_\infty(x_0, \epsilon)$  (left) is mapped by the neural network  $f$  to a region in the output space (right). Verification succeeds if the entire mapped region lies within the decision boundary (dotted red) of the correct class (Class A). This illustrates the robustness verification problem.

The earliest solutions that attempted to solve this problem are end-to-end/complete verification solutions. [4] introduced a tool called Reluplex, which is a solver based on SMT that combined the Simplex algorithm with ReLU activation functions to exactly verify small neural networks. Later, [7] used Mixed Integer Linear Programming (MILP) to write the problem as a commercial solution that exactly calculates robustness bounds. Although these solutions are sound and complete, they are NP-complete problems, known to not scale with the size of modern neural networks which are deployed in safety-critical applications.

For overcoming the scalability issue, Abstract Interpretation [1] and later for neural networks [2] has been adopted, which is less complete but more efficient because it over-approximates how the network behaves. [8] introduced a way to quickly compute robustness certificates with a linear bound propagation approach. The DeepPoly[6] approach pushed the boundaries even further by combining floating-point polyhedra with intervals, which, thanks to linearization, managed non-linearities such as Sigmoid and Tanh activations. More recent contributions, such as [5]

In our research, we ask the following question:

*Can we continue to prove robustness properties of a neural network by leveraging the properties of dual numbers and information about the gradients?*

We proceed with the foundation built on the *Abstract Dual Domain*. Our research is based on the theory of forward-mode automatic differentiation, which uses a concept called *dual numbers*[3]. We aim to analyze the sensitivity of the neural network to changes in the input, which can potentially provide a tighter bound on certain verification problems.

## 2 Overview

In this section, we introduce the key insight behind our approach and define the Abstract Dual Domain. Our approach is based on the following intuition: while relational abstractions such as DeepPoly[6] can be extremely precise, they often suffer from the “wrapping effect” where approximation errors accumulate layer by layer. In contrast, global gradient bounds can be sometimes a simpler yet effective verification domain, especially for shallow smooth networks.

### 2.1 Paradigm Shift: From Relational Bounds to Gradient Analysis

The basic underlying concept of this work is to effectively bridge the gap between standard Abstract Interpretation and Forward-Mode Automatic Differentiation. Standard methods like DeepPoly track affine constraints ( $x_j \geq \sum w_i x_i$ ) to bound the output values directly. However, precise reapproximating the feasible set at every non-linear layer introduces error which tends to accumulate with depth ( $O(L)$ ). Our approach, instead, exploits the Mean Value Theorem to find bounds on the output variation in terms of the gradient variation:

$$|f(x) - f(x_0)| \leq \underbrace{\sup_{z \in X} \|\nabla f(z)\|}_{\text{Abstract Duals}} \cdot \|x - x_0\| \quad (1)$$

By over-approximating the sound of the gradient  $\nabla f(z)$  for the entire input space  $X$ , we can guarantee robustness because the maximum possible change to the output will be too small to change the classification label. The key insight is that this “jumps” over the intermediate layers and thus avoids some of the wrapping errors from the layer-wise value propagation.

## 99    2.2 Abstract Dual Domain

100 To formalize this idea, we can start by lifting standard dual numbers to the domain of Affine  
 101 Arithmetic. Consider the Dual domain  $\hat{\mathcal{D}}$  as a product space of two affine forms representing the  
 102 range of values and the range of gradients across a set. Then we can say:

103    *Definition 2.1.* An **Abstract Dual Number**  $X \in \hat{\mathcal{D}}$  is a pair:

$$104 \quad X = \langle \hat{x}_{val}, \hat{x}_{grad} \rangle \quad (2)$$

107 where:

- 108    •  $\hat{x}_{val} = \alpha_0 + \sum_{i=1}^n \alpha_i \epsilon_i$  is the affine form representing the interval of neuron values.
- 109    •  $\hat{x}_{grad} = d_0 + \sum_{i=1}^n d_i \epsilon_i$  is the affine form representing the interval of partial derivatives with  
 110 respect to the input.

## 111    2.3 Propagation Rules

113 To propagate  $X$  through a neural network, we define abstract transformers for each layer type.

114    **Linear Layers:** For a fully connected layer with weight matrix  $W$  and bias vector  $b$ , the trans-  
 115 formation is exactly determined by the linearity of the dual algebra. The affine forms for value and  
 116 gradient are transformed as follows:

$$117 \quad Y = \langle W\hat{x}_{val} + b, W\hat{x}_{grad} \rangle \quad (3)$$

119 This operation is exact in the affine domain and introduces no new approximation errors.

120    **Non-linear Activations:** For a smooth non-linear activation function  $\sigma$  (e.g., Sigmoid), we  
 121 cannot directly apply the function to the affine forms. We handle the value and gradient components  
 122 separately:

123    *Value Component:* We approximate  $\sigma(\hat{x}_{val})$  using a linear relaxation. We linearize the function  
 124 around the center of the input affine form  $c$ . The new center is  $\sigma(c)$ , and the noise coefficients are  
 125 scaled by the derivative  $\sigma'(c)$ . To ensure soundness, we add a linearization error term to the radius  
 126  $r$ , which depends on the maximum curvature (second derivative) of  $\sigma$  and the radius of the input  
 127 interval.

$$128 \quad \hat{y}_{val} \approx \sigma(c) + \sigma'(c) \cdot (\hat{x}_{val} - c) + \epsilon_{err} \quad (4)$$

130    *Gradient Component:* By the chain rule, the gradient of the output is the product of the local  
 131 derivative and the input gradient:  $\nabla y = \sigma'(x) \cdot \nabla x$ . In our abstract domain, we compute an affine  
 132 approximation of the derivative  $\hat{\sigma}'$  based on the output value  $\hat{y}_{val}$ . For the Sigmoid function, we  
 133 use the property  $\sigma'(x) = \sigma(x)(1 - \sigma(x))$  to approximate the derivative as  $\hat{y}_{val} \otimes (1 - \hat{y}_{val})$ , where  
 134  $\otimes$  denotes affine multiplication. The output gradient is then obtained by multiplying this derivative  
 135 approximation with the input gradient affine form:

$$136 \quad \hat{y}_{grad} = \hat{\sigma}' \cdot \hat{x}_{grad} \quad (5)$$

138 This approach preserves the correlations between the derivative and the gradient, providing a more  
 139 precise abstraction than simple interval scaling.

## 140    2.4 Gradient Instability

142 A key insight from our analysis is the concept of *gradient instability*. We define a neuron  $n$  as  
 143 having an unstable gradient if its gradient interval  $[\underline{g}, \bar{g}]$  contains 0:

$$144 \quad \exists x_1, x_2 \in X_0 \text{ s.t. } \text{sign} \left( \frac{\partial f}{\partial n}(x_1) \right) \neq \text{sign} \left( \frac{\partial f}{\partial n}(x_2) \right) \quad (6)$$

148 Geometrically, this means that the function is non-monotonic w.r.t. neuron  $n$  over the input  
 149 region. When this happens, the assumption of local linearity breaks down and linear relaxations-  
 150 such as those in DeepPoly- become loose. Our approach explicitly tracks these gradient intervals,  
 151 thus enabling us to identify whenever and wherever the network's behaviour becomes hard to  
 152 certify linearly.

### 153 3 Contributions

#### 154 3.1 Theoretical Results: Soundness and Instability

#### 155 4 Formal Proof of Soundness

156 In this section, we establish the mathematical validity of the Abstract Dual domain. We demon-  
 157 strate that the interval of gradients produced by our forward-mode propagation soundly over-  
 158 approximates the true range of partial derivatives of the neural network  $f$  over the input region  
 159  $X_0$ .

#### 160 4.1 Foundational Definitions

161 Let  $f : \mathcal{R}^n \rightarrow \mathcal{R}^k$  be a neural network. Let  $X_0 \subseteq \mathcal{R}^n$  be a centrally symmetric input region  
 162 represented by the affine form  $\hat{x}_{val}$ . We define the **Concrete Derivative Set** as follows:

$$163 \quad \mathcal{D}(f, X_0) = \{\nabla f(x) \in \mathcal{R}^{k \times n} \mid x \in X_0\} \quad (7)$$

164 *Definition 4.1 (Soundness of Abstract Duals).* An Abstract Dual Number  $\mathcal{X} = \langle \hat{x}_{val}, \hat{x}_{grad} \rangle$  is  
 165 considered **sound** with respect to a function  $f$  and input region  $X_0$  if and only if:

- 166 (1)  $\forall x \in X_0 : f(x) \in \gamma(\hat{x}_{val})$
- 167 (2)  $\forall x \in X_0 : \nabla f(x) \in \gamma(\hat{x}_{grad})$

168 where  $\gamma(\hat{x})$  denotes the concretization function mapping an affine form to its corresponding subset  
 169 of  $\mathcal{R}$ .

#### 170 4.2 Linear Layer Soundness

171 *LEMMA 4.2 (SOUNDNESS OF LINEAR TRANSFORMERS).* Given a sound abstract dual  $\mathcal{X}$  and a linear  
 172 transformation  $y = Wx + b$ , the abstract transformer  $\mathcal{Y} = W\mathcal{X} + b$  is sound.

173 *PROOF.* By the fundamental rules of multi-variable calculus, if  $y$  is a linear function of  $x$  defined  
 174 by  $y = Wx + b$ , the Jacobian matrix is constant:  $\nabla_x y = W$ . By the chain rule, for a composite  
 175 function  $y(f(x))$ , we have:

$$176 \quad \nabla y = W \cdot \nabla f(x) \quad (8)$$

177 In our abstract domain, we define:

$$178 \quad \hat{y}_{val} = W\hat{x}_{val} + b, \quad \hat{y}_{grad} = W\hat{x}_{grad} \quad (9)$$

179 Since affine arithmetic is a linear abstraction, it is exact for linear transformations (i.e.,  $\gamma(W\hat{x} + b) =$   
 180  $\{Wx + b \mid x \in \gamma(\hat{x})\}$ ). Given that  $\mathcal{X}$  is sound,  $\gamma(\hat{x}_{grad})$  contains all possible values of  $\nabla f(x)$ .  
 181 Therefore,  $\gamma(W\hat{x}_{grad})$  must contain  $W \cdot \nabla f(x)$ . This concludes that  $\mathcal{Y}$  is sound.  $\square$

#### 182 4.3 Activation Function Soundness

183 *LEMMA 4.3 (SOUNDNESS OF NON-LINEAR ACTIVATION TRANSFORMERS).* Let  $\sigma$  be a differentiable  
 184 activation function (e.g., Sigmoid, Tanh). The abstract dual transformer  $\sigma^\#(\mathcal{X})$  defined by interval  
 185 derivative bounding is sound.

PROOF. Consider the composite function  $h(x) = \sigma(f(x))$ . By the univariate chain rule applied element-wise:

$$\nabla h(x) = \sigma'(f(x)) \cdot \nabla f(x) \quad (10)$$

Let  $[l, u]$  be the interval range of the values  $\hat{x}_{val}$ . We compute the derivative range  $\Sigma'$  as:

$$\Sigma' = \left[ \inf_{z \in [l, u]} \sigma'(z), \sup_{z \in [l, u]} \sigma'(z) \right] \quad (11)$$

The abstract gradient is then computed as the product  $\hat{h}_{grad} = \Sigma' \cdot \hat{x}_{grad}$ . By the Mean Value Theorem, for any  $x \in X_0$ , the concrete value  $f(x)$  must lie within  $[l, u]$ . Consequently, the concrete derivative  $\sigma'(f(x))$  must be contained within the interval  $\Sigma'$ . Using the soundness property of interval-affine multiplication, the resulting affine form  $\hat{h}_{grad}$  is guaranteed to enclose the product of any value in  $\Sigma'$  and any vector in  $\gamma(\hat{x}_{grad})$ . Thus,  $\forall x \in X_0 : \nabla h(x) \in \gamma(\hat{h}_{grad})$ .  $\square$

#### 4.4 Main Theorem: Lipschitz Certification

**THEOREM 4.4 (GLOBAL LIPSCHITZ SOUNDNESS).** *The Lipschitz constant  $K_{comp}$  derived from the output Abstract Dual  $\mathcal{Y}_{out}$  is a sound upper bound for the function  $f$  over the region  $X_0$ .*

PROOF. By induction over the network depth  $L$ , using Lemmas 1 and 2, the final gradient affine form  $\hat{y}_{grad}$  at the output layer is a sound over-approximation of the set  $\mathcal{D}(f, X_0)$ . The global Lipschitz constant for  $f$  with respect to the  $L_\infty$  norm is defined as:

$$K = \sup_{x \in X_0} \|\nabla f(x)\|_1 \quad (12)$$

For any affine form  $\hat{g} = \alpha_0 + \sum_{i=1}^n \alpha_i \epsilon_i$ , the supremum of its absolute value is soundly bounded by:

$$\sup |\gamma(\hat{g})| \leq |\alpha_0| + \sum_{i=1}^n |\alpha_i| \quad (13)$$

We define  $K_{comp}$  as the sum of these absolute coefficients for the output gradient affine form. It follows that  $K_{comp} \geq K$ . By the Mean Value Theorem:

$$\forall x \in X_0 : \|f(x) - f(x_0)\|_\infty \leq K_{comp} \cdot \|x - x_0\|_\infty \quad (14)$$

Thus, if the safety condition  $f(x_0)_{target} - f(x_0)_{other} > K_{comp} \cdot \epsilon$  is satisfied, the classification is guaranteed to be invariant within the  $\epsilon$ -ball.  $\square$

#### Acknowledgments

TBD

#### References

- [1] Patrick Cousot and Radhia Cousot. 1977. Abstract interpretation. *Symposium on Principles of Programming Languages* (Jan 1977). doi:10.1145/512950.512973
- [2] Timon Gehr, Matthew Mirman, Dana Drachsler-Cohen, Petar Tsankov, Swarat Chaudhuri, and Martin T. Vechev. 2018. AI2: Safety and Robustness Certification of Neural Networks with Abstract Interpretation. *2018 IEEE Symposium on Security and Privacy (SP)* (2018), 3–18. <https://api.semanticscholar.org/CorpusID:206579396>
- [3] Andreas Griewank and Andrea Walther. 2008. *Evaluating derivatives: principles and techniques of algorithmic differentiation*. SIAM.
- [4] Guy Katz, Clark Barrett, David L Dill, Kyle Julian, and Mykel J Kochenderfer. 2017. Reluplex: An efficient SMT solver for verifying deep neural networks. In *International Conference on Computer Aided Verification*. Springer, 97–117.
- [5] Gagandeep Singh, Timon Gehr, Matthew Mirman, Markus Püschel, and Martin Vechev. 2018. Fast and Effective Robustness Certification. In *Advances in Neural Information Processing Systems*, S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett (Eds.), Vol. 31. Curran Associates, Inc. [https://proceedings.neurips.cc/paper\\_files/paper/2018/file/f2f446980d8e971ef3da97af089481c3-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2018/file/f2f446980d8e971ef3da97af089481c3-Paper.pdf)

- 246 [6] Gagandeep Singh, Timon Gehr, Markus Püschel, and Martin Vechev. 2019. An abstract domain for certifying neural  
247 networks, Vol. 3. 1–30. doi:10.1145/3290354
- 248 [7] Vincent Tjeng, Kai Xiao, and Russ Tedrake. 2019. Evaluating Robustness of Neural Networks with Mixed Integer  
249 Programming. In *International Conference on Learning Representations*.
- 250 [8] Tsui-Wei Weng, Huan Zhang, Hongge Chen, Zhao Song, Cho-Jui Hsieh, Duane S. Boning, Inderjit S. Dhillon, and Luca  
251 Daniel. 2018. Towards Fast Computation of Certified Robustness for ReLU Networks. *ArXiv* abs/1804.09699 (2018).  
252 <https://api.semanticscholar.org/CorpusID:13750928>
- 253
- 254
- 255
- 256
- 257
- 258
- 259
- 260
- 261
- 262
- 263
- 264
- 265
- 266
- 267
- 268
- 269
- 270
- 271
- 272
- 273
- 274
- 275
- 276
- 277
- 278
- 279
- 280
- 281
- 282
- 283
- 284
- 285
- 286
- 287
- 288
- 289
- 290
- 291
- 292
- 293
- 294