

Certifying Differential Invariants of Neural Networks using Abstract Duals

CHANDRA KANTH NAGESH

*Department of Computer Science
University of Colorado Boulder*

Final Project; CSCI – 7135, Fall 2025



Safety Property

Given a neural network classifier $f : \mathbb{R}^n \rightarrow \mathbb{R}^k$, an input x_0 , and a perturbation radius ϵ . Verify that the classification of x_0 is invariant within the ϵ -ball:

$$\forall x \in \mathbb{B}_\infty(x_0, \epsilon) : \quad \text{argmax}(f(x)) = \text{argmax}(f(x_0))$$



Original Image

Original Image
 $+ \epsilon = 0.25$

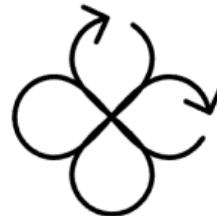
Original Image
 $+ \epsilon = 0.1$

The Problem

Neural Networks are highly non-linear and vulnerable to *adversarial attacks*.



Original Image + tiny perturbations



Continuous input spaces



Safety-Critical Domains

Possible solution

Try computing an over-approximation of the reachable output set $\mathcal{R}_f(X_0)$ to prove safety

Why is it Interesting?

The power of global bounds with simplicity of gradient analysis

1. Simplicity beats complexity

- DeepPoly¹ is strictly better than our method.
- However, *surprising results found!*

2. Gradient Instability

- Found 100% correlation between instability and failure at high ϵ
- Our method detects when the *local linearity assumption collapses*

$$f(x)_{true} - f(x)_{other} > \sup_{x \in X_0} \|\nabla f(x)\|_1 \cdot \epsilon$$

¹Singh, G. et. al. (2019). An abstract domain for certifying neural networks. 3(POPL), pp.1–30.

Why is it Hard?

1. Activation Non-Linearity:

- **Sigmoid ($\sigma(x)$)**: Non-convex. Linear relaxation introduce loose bounding boxes
- **ReLU ($\max(0, x)$)**: Discontinuous derivative
 - ▶ $\nabla \text{ReLU}(x) \in \{0, 1\}$. At $x = 0$, the sub-gradient is the set $[0, 1]$
 - ▶ This causes the Lipschitz bound to explode: $K_{layer} = \|W\| \cdot 1$

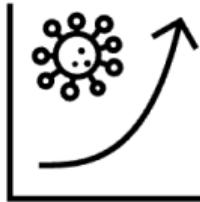
2. The Wrapping Effect:

- Transforming a simple shape (zonotope/box) through a linear map rotates it
- Re-approximating it with axis-aligned bounds adds "dead space" volume exponentially with depth

Contributions



Formalization:
Affine Arithmetic



Gradient Anomaly:
Depth < 3; Sigmoid



Instability Metric:
Identifies failures



Evaluation:
Tiny/Small/Std. ϵ

Statement: "Global gradient bounds provide a superior verification domain for shallow, smooth networks, while relational domains are essential for deep, non-smooth architectures"

Contribution Idea

Algebra $\mathbb{D} = \{a + b\eta\}$. Exact derivative propagation looks like: $f(x + \eta) = f(x) + f'(x)\eta$

- We lift this to **Affine Forms** (Zonotopes) to represent the input region X_0 :
- $\hat{x} = \alpha_0 + \sum_{i=1}^n \alpha_i \epsilon_i, \quad \epsilon_i \in [-1, 1]$
- Forward Propagation:

$$\hat{y} = f(\hat{x}) = \underbrace{\left(y_0 + \sum y_i \epsilon_i \right)}_{\text{Value Interval}} + \underbrace{\left(d_0 + \sum d_i \epsilon_i \right)}_{\text{Gradient } \nabla f} \eta$$

- Lipschitz Bound:

$$K = \sup_{x \in X_0} ||\nabla f(x)|| = |d_0| + \sum |d_i|$$

Contribution Idea

A neuron n is **gradient-unstable** if its gradient interval $[\underline{g}, \bar{g}]$ contains 0

$$\exists x_1, x_2 \in X_0 \text{ s.t. } \operatorname{sign} \left(\frac{\partial f}{\partial n}(x_1) \right) \neq \operatorname{sign} \left(\frac{\partial f}{\partial n}(x_2) \right)$$

Implication

- Instability \implies High local curvature (non-monotonicity)
- Linear approximations (DeepPoly) become loose when curvature is high

1. **Sigmoid (Smooth):**
 - Gradient method wins in **16.1%** of Tiny nets.
 - *Reason:* Smooth derivatives allow tight global bounds via Abstract Duals.
2. **ReLU (Non-smooth):**
 - DeepPoly dominates (**100% win rate**).
 - *Reason:* Discontinuous gradients ($\{0, 1\}$) cause the abstract dual interval to become $[0, 1]$ everywhere, destroying precision.
3. **Performance:** Gradient analysis is **2.4x faster** but scales poorly with depth ($K_{total} \approx \prod ||W_i||$).

OCaml code run on MNIST classifiers (Tiny: 2x10, Small: 5x20, Standard: 10x20) with perturbation $\epsilon \in \{0.01 \dots 0.12\}$.

- No “Silver Bullet” domain exists for Neural Network verification
- **Shallow + Smooth** → Use **Gradient (Abstract Duals)**
- **Deep + Non-smooth** → Use **Relational (DeepPoly)**

Final Takeaway

“Precision in static analysis is not just about the domain’s complexity, but its alignment with the function’s smoothness.”